

1. Imports and Project Environment

- I have turned this Notebook file into a multipurpose project report because combining the Markdown and Python output was very convenient.
- Some functions that I have written are inside the `/project/util.py` file to reduce visual clutter.

```
In [3]: from project.util import *
```

```
In [4]: %watermark -v -p sklearn,pandas,numpy,matplotlib,seaborn
```

```
Python implementation: CPython
Python version      : 3.13.7
IPython version    : 9.6.0

sklearn   : 1.7.2
pandas    : 2.3.3
numpy     : 2.3.4
matplotlib: 3.10.7
seaborn   : 0.13.2
```

2. Data Preparation

- Reading the raw data in CSV format and renaming the columns
- Extracting labels (Diagnosis) from the features and mapping classes as "Malignant -> 1" and "Benign -> 0"
- Normalizing features using MinMaxScaler

```
# Column names taken from the "Variables Table" part of the dataset source
col_names = ['ID', 'Diagnosis',
             'radius1', 'texture1', 'perimeter1', 'area1', 'smoothness1', 'compactness1',
             'concavity1', 'concave_points1', 'symmetry1', 'fractal_dimension1',
             'radius2', 'texture2', 'perimeter2', 'area2', 'smoothness2', 'compactness2',
             'concavity2', 'concave_points2', 'symmetry2', 'fractal_dimension2',
             'radius3', 'texture3', 'perimeter3', 'area3', 'smoothness3', 'compactness3',
             'concavity3', 'concave_points3', 'symmetry3', 'fractal_dimension3']

data_df = pd.read_csv('data/wdbc.data', sep=',', names=col_names)
display_pruned_df(data_df)
```

```
Out[5]:
```

	ID	Diagnosis	radius1	texture1	...	concave_points3	symmetry3	fractal_dimension3
0	842302	M	17.99	10.38	...	0.2654	0.4601	0.1189
1	842517	M	20.57	17.77	...	0.186	0.275	0.08902
...
567	927241	M	20.6	29.33	...	0.265	0.4087	0.124
568	92751	B	7.76	24.54	...	0.0	0.2871	0.07039

569 rows x 32 columns

```
In [6]: features = data_df.drop(['ID'], axis=1)
labels = features.pop('Diagnosis').map({'B': 0, 'M': 1})
```

```
In [7]: scaler = MinMaxScaler()
features = pd.DataFrame(scaler.fit_transform(features), columns=col_names[2:])
display_pruned_df(features)
```

Out[7]:


	radius1	texture1	perimeter1	area1	...	concave_points3	symmetry3	fractal_dimension3
0	0.521037	0.022658	0.545989	0.363733	...	0.912027	0.598462	0.418864
1	0.643144	0.272574	0.615783	0.501591	...	0.639175	0.23359	0.222878
...
567	0.644564	0.66351	0.665538	0.475716	...	0.910653	0.497142	0.452315
568	0.036869	0.501522	0.02854	0.015907	...	0.0	0.257441	0.100682

569 rows x 30 columns

3. Distribution of Features

In [8]:

```
for i in range(3):
    f, ax = plt.subplots(figsize=(14, 4))
    subset = features.iloc[:, i*10:(i+1)*10].melt(var_name='Feature', value_name='Value')
    sns.boxenplot(data=subset, x='Feature', y='Value', ax=ax)
    plt.tight_layout()
    plt.show()
```




4. Determining the Number of Clusters

- The elbow point can be selected as `k=2`, since the change in inertia becomes less significant after that point.
- While determining the number of clusters, higher silhouette score is desired. So, `k=2` is once again the best choice.

- Therefore, we can safely assume that the number of clusters is 2.

```
In [9]: f, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 4))
```

```
plot_elbow_method_with_k(ax1, features, k_start=1, k_end=9)
plot_silhouette_score_with_k(ax2, features, k_start=2, k_end=9)
```



5. Clustering and Evaluation

5.1. K-Means

- `k=2` was found as the number of clusters in the previous section, so we use it here as the parameter `n_clusters=2`.
- Using Adjusted Rand Index (ARI) as the evaluation metric (which ranges between 0 and 1) gives a decent result.


```
In [10]: k_means = KMeans(n_clusters=2, random_state=0).fit(features)
k_means_ari_score = adjusted_rand_score(labels, k_means.labels_)
print(f'+ K-Means ARI Score: {k_means_ari_score:.6f}')
```

```
+ K-Means ARI Score: 0.730175
```


5.2. DBSCAN

- First, I have used the range `2 <= minPts <= 11` to find the optimal eps ranges using 10th and 95th percentile distances of Nearest Neighbor pairs.
- When evaluated with the Silhouette Score metric, the scores for `minPts=[2,3]` seems ill-behaved even after the potential knee-point. Therefore, I have decided to use `minPts=4` going forward.
- Furthermore, evaluating the DBSCAN labels against the ground truth using ARI metric yields a significantly worse score compared to the K-Means algorithm.

```
In [11]: dbSCAN_params_df = find_optimal_dbSCAN_params(features, min_samples_range=(2, 11))
```



```
In [12]: min_samples = 4
eps = find_optimal_eps(dbSCAN_params_df, min_samples=min_samples, min_eps=0.34)
```



```
In [13]: dbSCAN = DBSCAN(min_samples=min_samples, eps=eps)
dbSCAN_labels = dbSCAN.fit_predict(features)
dbSCAN_ari_score = adjusted_rand_score(labels, dbSCAN_labels)
print(f'+ DBSCAN ARI Score: {dbSCAN_ari_score:.4f}')
```


+ DBSCAN ARI Score: 0.2115

6. Visualisation on 2D Space

- To visualise the clusters on 2D, I have used PCA to get the two best components that represent the 30-dimensional space.

```
In [14]: f, (ax1, ax2, ax3) = plt.subplots(1, 3, sharey=True, figsize=(14, 5))
```

```
plot_on_pca(ax1, features, labels, title='Real Clusters')
plot_on_pca(ax2, features, k_means.labels_, title='K-Means Clusters (k=2)')
plot_on_pca(ax3, features, dbSCAN_labels, title=f'DBSCAN Clusters (minPts={min_samples} eps={eps})')
plt.tight_layout()
```



Ege Altıok - 2025900151